

# Machine Learning



**RAJESH SHARMA**

Walt Disney Animation Studios



Thank you to ACM SIGGRAPH!



**Pol Jeremias-Vila** : SIGGRAPH 2021 Chair

**Tomasz Bednarz**: Frontiers Program Chair

**Alex Bryant**: Student Volunteers Chair

**Tim Hendrickson**: Digital Marketing Manager

Student Volunteers:

**Rogelio, Trinity, Aurora, Emily, Hunter & Kendra**



**SIGGRAPH 2021**

# Machine Learning

————— Rajesh Sharma —————

# Marios Papas



## Research Scientist

### **Research Scientist**

Disney Research | Studios - Zurich

As the Rendering Group lead at Disney Research, Marios works with a talented team of researchers, engineers and students from Walt Disney Animation Studios, Pixar, Industrial Light and Magic and ETH Zurich.

His research focuses on Monte Carlo rendering problems to understand how light interacts with materials and use that knowledge to design state-of-the-art algorithms and appearance models for efficient and realistic image synthesis.

His current research focuses on using Machine Learning methods for information sharing between samples generated by path tracing algorithms for driving future sampling decisions (adaptive importance sampling) and achieve low-error reconstruction (denoising).

# Today

- Quick Recap: Regression, HW
- Second Neural Network - Classification
- Least Squares  $\rightarrow$  Maximum Likelihood
- Autoencoder

# Hands-on

- ★ Log in to your google drive
- ★ Make a shortcut to: `https://bit.ly/3oKCVCh`
- ★ Make a copy of:
  - Sinx.ipynb (Homework from last time)
  - FlowerClassification.ipynb
  - Autoencoder.ipynb

# Recap - Questions

Aidan Whelan - What kinds of models do you find are most commonly used in your work in the film industry?

Avidsiman - How do we know if the model has memorized our data instead of fitted it?

Anjoe Jacob - Can you also expand a bit about factors that go into choosing between mean avg err vs mean squared err?

# Classification


Given some pre-classified data:

Q: How do we predict which class the new data belongs to?

A: By finding the maximum likelihood




# Probability & Statistics: Normal Distribution



$N(\mu, \sigma^2)$  ;  $\mu$  - mean ,  $\sigma^2$  - variance

$$f(x | \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \cdot e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$


# Probability & Statistics: Normal Distribution



$N(\mu, \sigma^2)$  ;  $\mu$  - mean ,  $\sigma^2$  - variance

$$f(x | \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \cdot e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

# Regression → Maximum Likelihood Estimation



Prediction:  $\hat{y} = ax + b$

Actual:  $y_i$


Error:  $|y_i - \hat{y}_i|$

Total Squared Error:  
 $\sum (y_i - \hat{y}_i)^2$ , for  $i=(1, n)$

Minimize Total Squared Error:  
 $E(a, b) = \sum (y_i - ax_i - b)^2$

$(a, b)$  are the parameters (weights)


# Regression → Maximum Likelihood Estimation



Prediction:  $\hat{y} = ax + b$

Actual:  $y_i$

Error:  $|y_i - \hat{y}_i|$



# Maximum Likelihood Estimation vs MSE


$$f(x \mid \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \cdot e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

$$f(x_1, \dots, x_n \mid \mu, \sigma^2) = \prod_{i=1}^n f(x_i \mid \mu, \sigma^2) = \left( \frac{1}{2\pi\sigma^2} \right)^{n/2} \exp\left( -\frac{\sum_{i=1}^n (x_i - \mu)^2}{2\sigma^2} \right).$$


$$\log \left( \mathcal{L}(\mu, \sigma) \right) = -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2$$

# Autoencoder

For regression, we had a fully-connected network, output layer size=1



# Autoencoder




# Autoencoder

Original Input

Latent Representation

Reconstructed Output





# Hands-on

- ★ *Log in to your google drive*
- ★ Find the shared folder 'Disney Machine Learning Webinar'
- ★ Make a copy of:
  - AutoEncoder.ipynb

# Autoencoder - Model


```
# build an autoencoder
model = tf.keras.models.Sequential([
    tf.keras.layers.InputLayer(IMG_SHAPE),
    tf.keras.layers.Flatten(),
    # encoder
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(32, activation='relu'),
    # decoder
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(784, activation='sigmoid'),
    tf.keras.layers.Reshape(IMG_SHAPE)
])

# compile
model.compile(optimizer='adamax', loss='mse')

# fit
model.fit(x_train, x_train, epochs=17, batch_size=256, shuffle=True, validation_data=(x_test, x_test))

# predict
decoded_imgs = model.predict(x_test)
```

# Autoencoder - results



Compression Factor:  $28 \times 28 / 32 \sim 25X$

# Caution

Remarkably Clever

Surprisingly Dumb



# Next Class

- Efficient Data Pipeline
- Convolutional Neural Network
- Artistic Style Transfer
- Homework:
  - Use the mnist dataset for classification
  - Extra credit: also show “next likely”
- @xarmalarma, #siggraph2021