

Machine Learning



RAJESH SHARMA

Walt Disney Animation Studios



Thank you to ACM SIGGRAPH!



Pol Jeremias-Vila : SIGGRAPH 2021 Chair

Tomasz Bednarz: Frontiers Program Chair

Alex Bryant: Student Volunteers Chair

Tim Hendrickson: Digital Marketing Manager

Student Volunteers:

Rogelio, Trinity, Aurora, Emily, Hunter & Kendra




SIGGRAPH 2021

Machine Learning

————— Rajesh Sharma —————

Leon Gatys




Research Scientist Apple

Leon Gatys is a Machine Learning researcher with a focus on modeling human experience. In his PhD thesis, he invented the popular [Neural Style Transfer](#) algorithm by using Deep Neural Networks to model Visual Perception. He is currently based in Seattle where he works as a founding member of Apple's Health AI team.

Research Scientist

Artistic Style Transfer



$$L_{\text{total}}(\check{c}, \check{s}, \check{x}) = \alpha L_{\text{content}}(\check{c}, \check{x}) + \beta L_{\text{style}}(\check{s}, \check{x})$$

Today


- Quick Recap: Distributions, Autoencoder
- Convolutional Neural Networks
- DataPipeline
- Denoising
- Transfer Learning / PreTrained Models

Hands-on

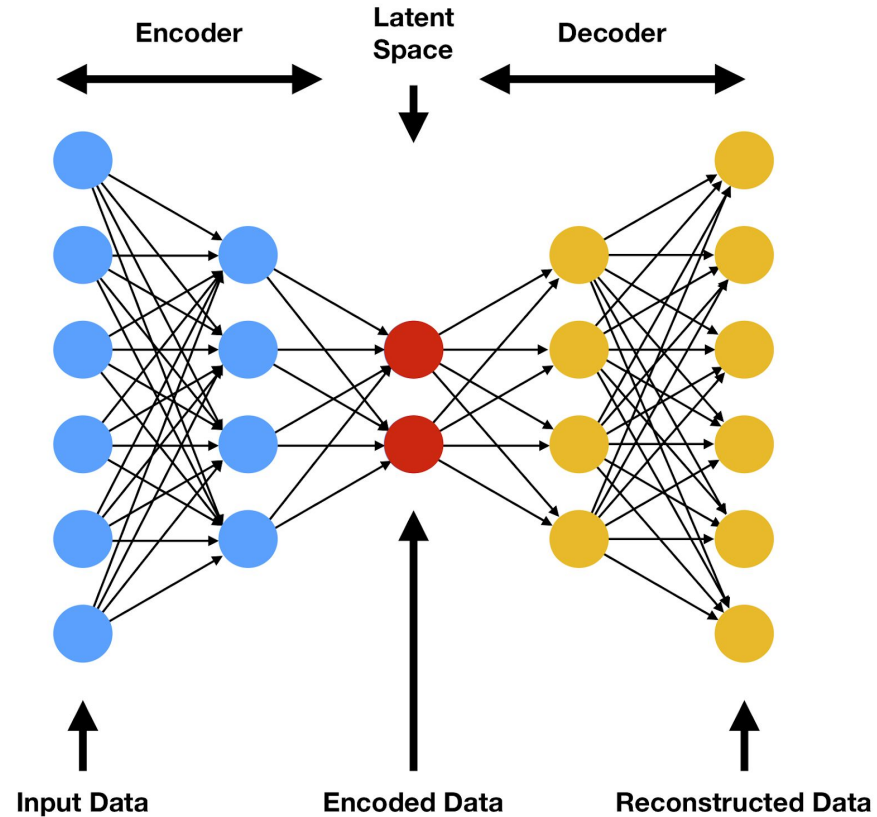
- ★ Log in to your google drive
- ★ Make a shortcut to: `https://bit.ly/3oKCVCh`
- ★ Make a copy of:
 - `Autoencoder.ipynb`
 - `dataPipeline.ipynb`
 - `denoiserCNN.ipynb`
 - `styleTransfer.ipynb`
 - `facialRecognition01.ipynb`

Autoencoder

For regression, we had a fully-connected network, output layer size=1



Autoencoder

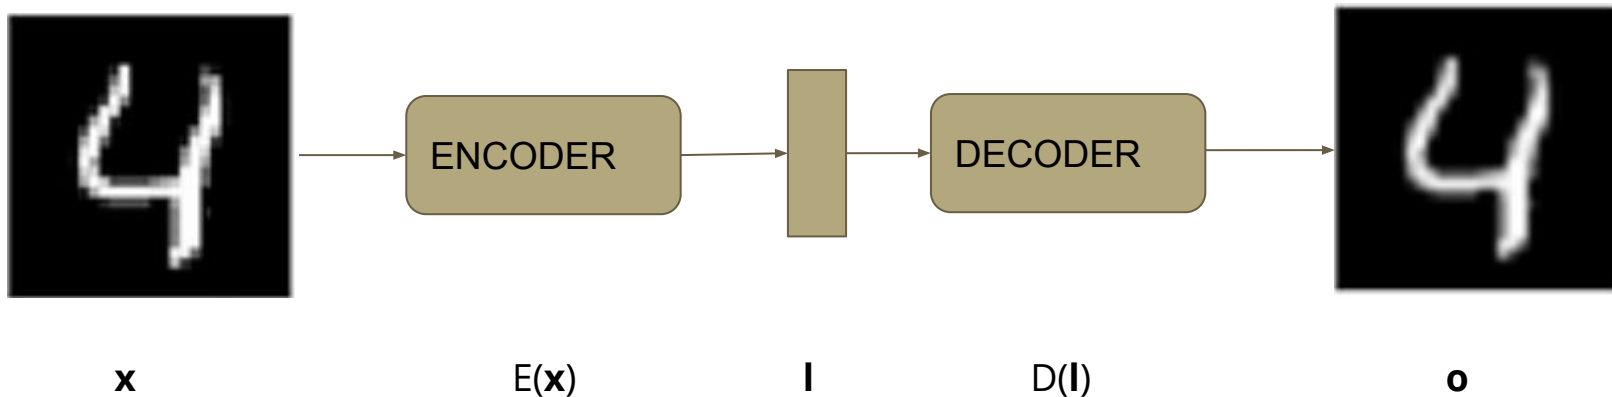


Autoencoder

Original Input

Latent Representation

Reconstructed Output



Autoencoder - Model

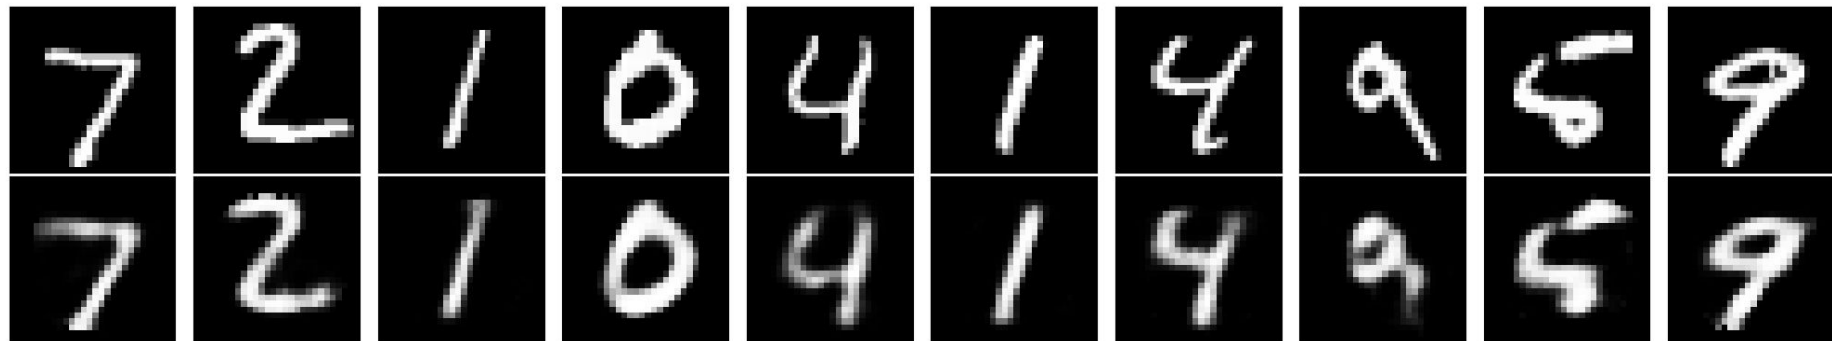
```
# build an autoencoder
model = tf.keras.models.Sequential([
    tf.keras.layers.InputLayer(IMG_SHAPE),
    tf.keras.layers.Flatten(),
    # encoder
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(32, activation='relu'),
    # decoder
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(784, activation='sigmoid'),
    tf.keras.layers.Reshape(IMG_SHAPE)
])

# compile
model.compile(optimizer='adamax', loss='mse')

# fit
model.fit(x_train, x_train, epochs=17, batch_size=256, shuffle=True, validation_data=(x_test, x_test))

# predict
decoded_imgs = model.predict(x_test)
```

Autoencoder - results



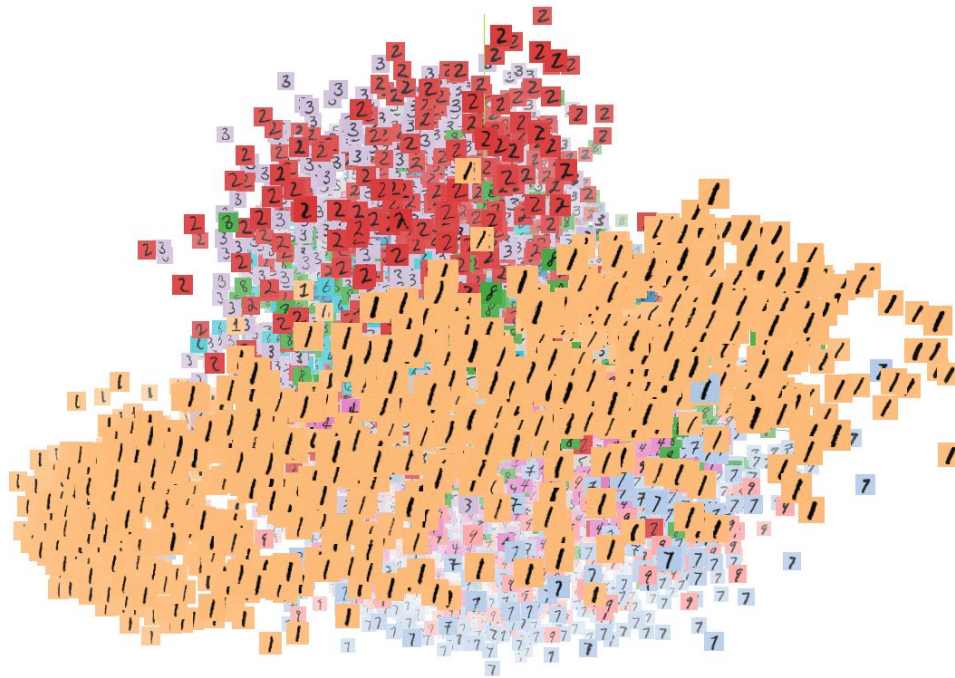
Compression Factor: $28 \times 28 / 32 \sim 25X$

Hands-on

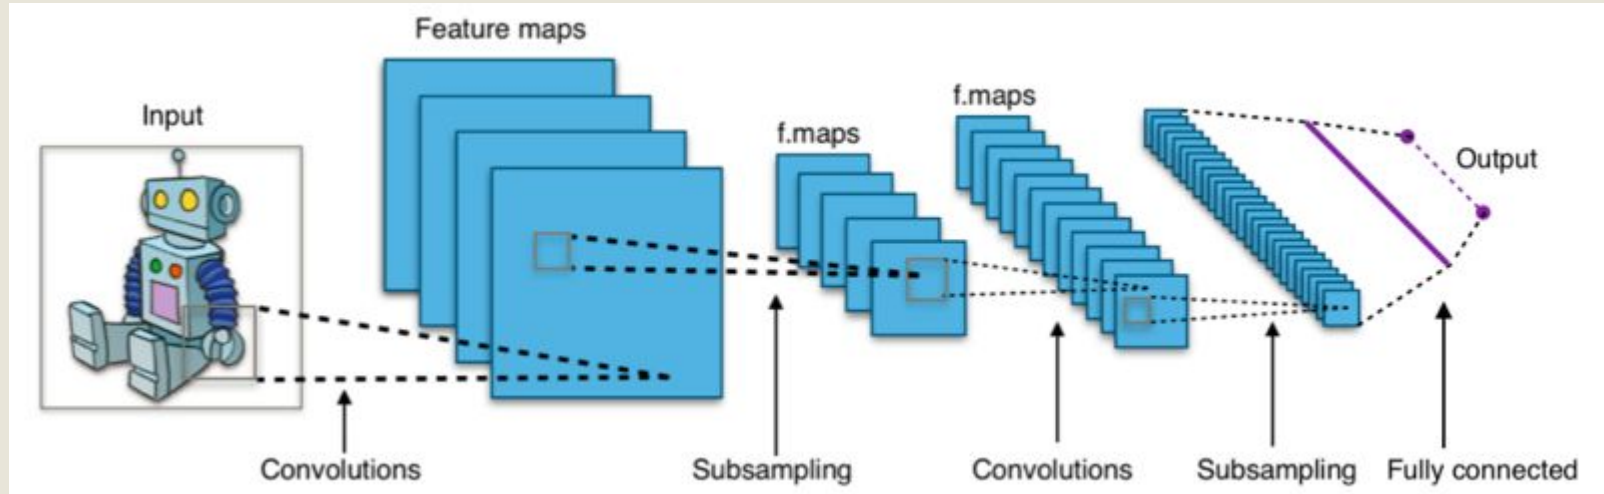
- ★ *Log in to your google drive*
- ★ Find the shared folder
- ★ Make a copy of:
 - AutoEncoder.ipynb

Latent Spaces and Embeddings

<https://projector.tensorflow.org>



Convolutional Neural Network (CNN)



Convolution (Extract High-Level Features)

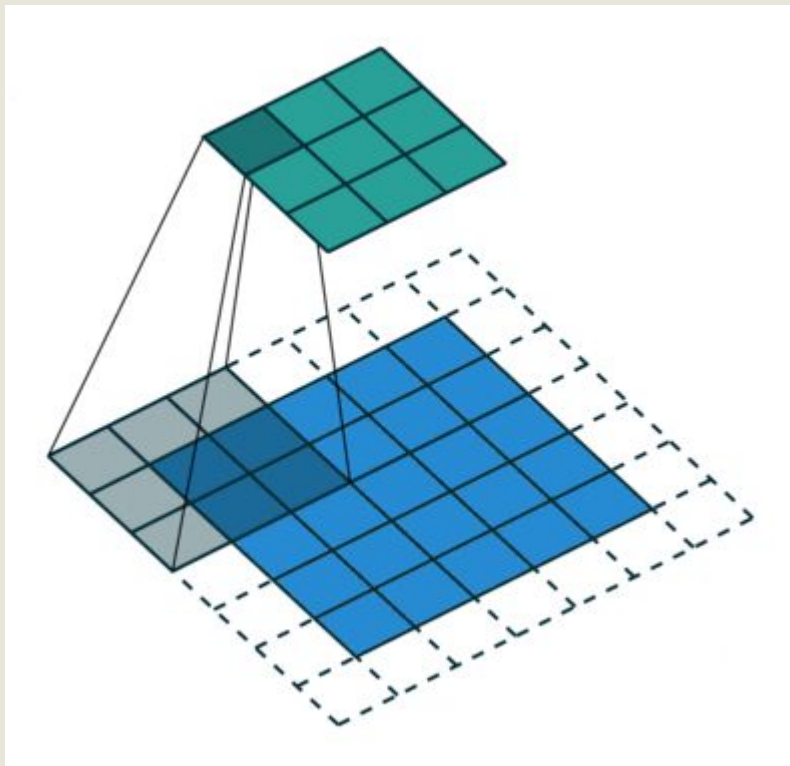



Image Denoising



Denoising with Kernel Prediction and Asymmetric Loss Functions SIGGRAPH 2018, Vogels et al

Noisy image....<similar image>....Clean image



Denoising with Kernel Prediction and Asymmetric Loss Functions SIGGRAPH 2018, Vogels et al

Noisy image.....<similar image>.....Clean image

- If we have a set of noisy images and, a set of corresponding clean images,
- We can train our network to recover
 - Clean images from noisy images
- How
 - By setting Clean image as the ground truth,
 - the Noisy image as input and,
 - the loss function as the difference btwn the two

Don't have a noisy version?

- Take a clean image
- Add synthetic noise to it (Data Augmentation)

But first, we need some more Engineering!

- Take a look at `dataPipeline.ipynb`
 - `--tensorflow` data sets and pipeline
 - `--addNoise`
 - `--extractPatches`

Noisy image...<similar image>...Clean image

- Take a look at denoiserCNN.ipynb
 - Make a CNN

—Noisy image.....<similar image>.....Clean image

degraded

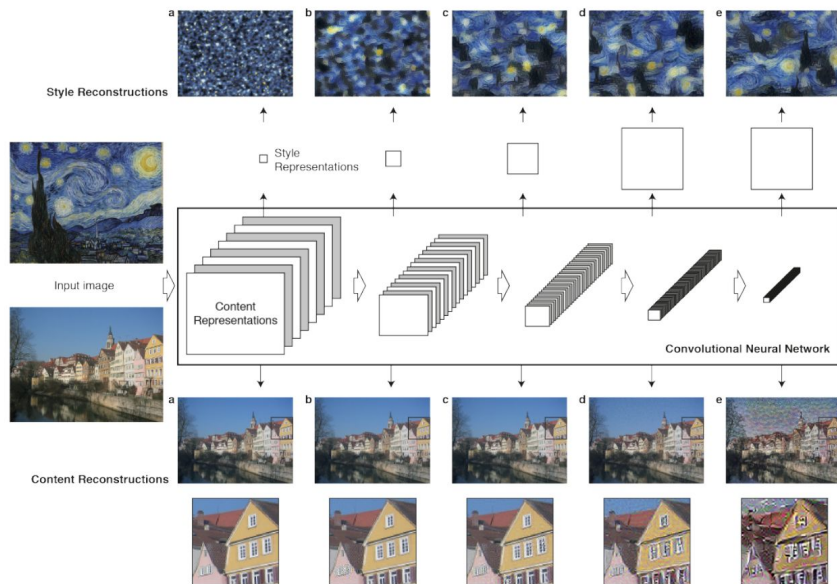
What else can we do?

- | | |
|---------------|---|
| Tint removal: | Image with tint.....<similar image>.....Clean image |
| In-painting: | Image with holes....<similar image>...Clean image |
| Dirt-removal: | Image with speckle....<similar image>...Clean image |
| Colorization: | Grayscale Image....<similar image>....Color image |
| Up-resing: | Lowres Image....<similar image>....Hires image |
| Inbetweening: | Image1, Image3....<similar image>.... Image2 |

Using off-the-shelf pretrained models



- Style Transfer
- MT-CNN

Artistic Style Transfer



$$L_{\text{total}}(\check{c}, \check{s}, \check{x}) = \alpha L_{\text{content}}(\check{c}, \check{x}) + \beta L_{\text{style}}(\check{s}, \check{x})$$

Extracting Faces -- MT-CNN



Homework:

Colorization: `tf.image.adjust_saturation`

Up-resing:

```
tf.image.resize(image, size=[256,256], method=tf.image.ResizeMethod.NEAREST_NEIGHBOR)
```

In-Painting:

```
mask = np.ones((PATCH_WIDTH, PATCH_HEIGHT), dtype=np.float32)
scale = 0.25
low, upper = int(PATCH_WIDTH * scale), int(PATCH_HEIGHT * (1.0 - scale))
mask[:, low:upper, low:upper] = 0.
tf.multiply(patch, mask)
```

Frame interpolation:

```
stacked = tf.concat([frame1, frame3], axis=-1)
```

Next Class

- Generative Neural Networks:
 - Variational AutoEncoder
 - Generative Adversarial Networks
- Homework:
 - Do other kinds of 'denoising'
- @xarmalarma, #siggraph2021